



# Cloud Computing

# Jellyfish in Data Centers

Aditi Sen & Charissa Zou



## Paper

- Paper: <http://pbg.cs.illinois.edu/papers/jellyfish-nsdi12.pdf>

### Jellyfish: Networking Data Centers Randomly

Ankit Singla<sup>†,\*</sup>, Chi-Yao Hong<sup>†,\*</sup>, Lucian Popa<sup>#</sup>, P. Brighten Godfrey<sup>†</sup>

<sup>†</sup> University of Illinois at Urbana–Champaign


<sup>#</sup> HP Labs

---

# Our Project

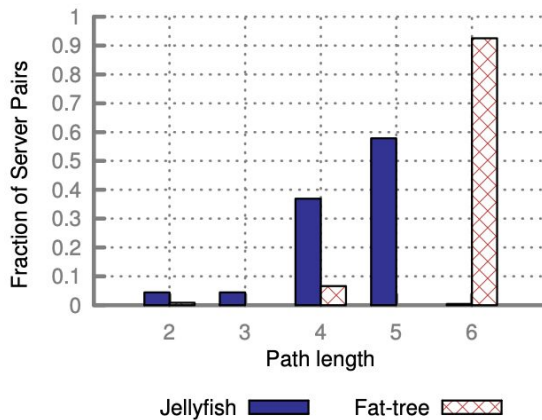

# Goals

1. Implement a degree bounded random graph Jellyfish network to simulate data center traffic



# Goals

2. Reproduce Figure 9 in the paper
  - k-shortest path provides better path diversity than ECMP for Jellyfish
3. Reproduce Figure 1c in the paper
  - Jellyfish produces shorter paths on average





## Reach Goal

4. Reproduce Table 1 in the paper
- Compare Jellyfish and Fat Tree throughput under different routing protocols
  - TCP vs. MPTCP
  - With the same switch equipment, Jellyfish supports 25% more servers than fat tree

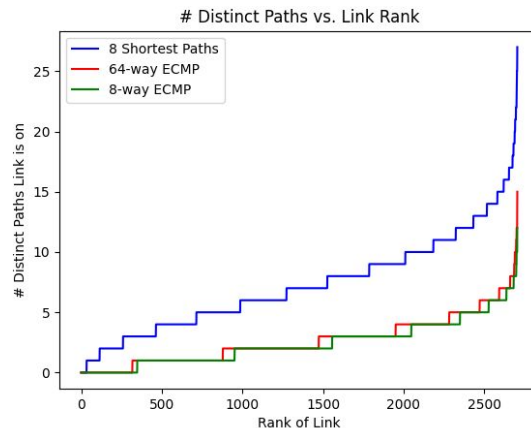

Congestion control	Fat-tree (686 svrs)	Jellyfish (780 svrs)	
	ECMP	ECMP	8-shortest paths
TCP 1 flow	48.0%	57.9%	48.3%
TCP 8 flows	92.2%	73.9%	92.3%
MPTCP 8 subflows	93.6%	76.4%	95.1%



Progress

# Last Time


- Implemented a Jellyfish random topology using Networkx in Python
- Simulated “traffic” direction and computed ECMP and k-Shortest Paths



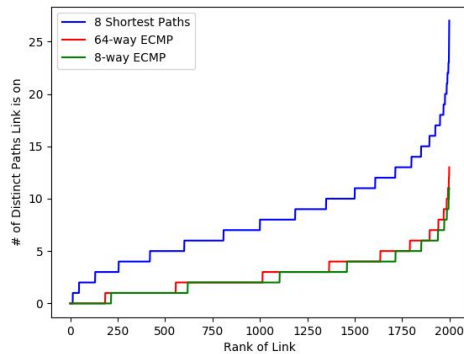


# Experiments


- Ran more experiments with a different number of switches and hosts



100 switches, 10 links




200 switches, 10 links




300 switches, 10 links

# Experiments

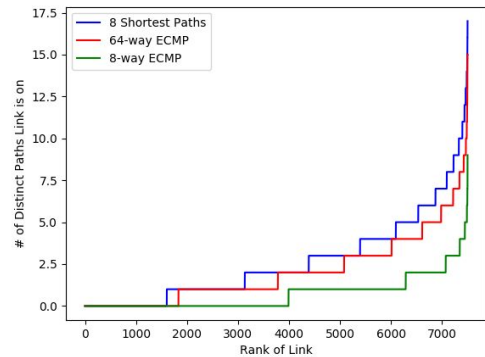
- Ran more experiments with a different number of switches and hosts



300 switches, 5 links




300 switches, 15 links



300 switches, 25 links

## Reproducing Table 1



Congestion control	Fat-tree (686 svrs)		Jellyfish (780 svrs)	
	ECMP		ECMP	8-shortest paths
TCP 1 flow	48.0%		57.9%	48.3%
TCP 8 flows	92.2%		73.9%	92.3%
MPTCP 8 subflows	93.6%		76.4%	95.1%



## Reproducing Table 1

- Built a Mininet jellyfish topology from Networkx graph

```
# Construct mininet
for n in graph.nodes:
    net.addSwitch("s_%s" % n)

    # Add single host on designated switches
    if int(n) in list(range(host_range)):
        net.addHost("h%s" % n)
        # directly add the link between hosts and their gateways
        net.addLink("s_%s" % n, "h%s" % n)

# Connect switches to each other as defined in networkx graph
for (n1, n2) in graph.edges:
    net.addLink('s_%s' % n1, 's_%s' % n2)
```



## Reproducing Table 1

- Tested on Mininet with a pre-built Pox controller
  - forwarding.l2\_learning

```
mininet> h0 ping h1
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data:
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=7.41 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.888 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.888 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.888 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.888 ms
^C

```

```
--- 10.0.0.2 ping statistics ---
376 packets transmitted, 376 received, 0% packet loss, time 383574ms
rtt min/avg/max/mdev = 0.040/0.187/7.407/0.673 ms
mininet>
```



## Reproducing Table 1

- Custom topology for Fat Tree
- Implement custom controllers with ECMP and 8-Shortest Paths routing
  - Pox + ripl vs. Ryu + OpenFlow
- Generate traffic flow using pre-constructed traffic matrix
- Measure throughput using iperf
- Parse and generate graphs based on data output

---

# Conclusion



# Challenges

- riplpox (Brandon Heller)
  - Version mismatch / permissions issues
- Pivoted to building everything from scratch
  - Mininet permission issues
  - Approach decisions
  - Custom controllers





# Looking Forward

- Complete Fat Tree topology construction (for comparison to Jellyfish)
- Implement ECMP and K-Shortest Path routing in custom controllers
- Simulate network traffic in Mininet
  - Measure throughput with iperf
- Generate graphs for network throughput
- Generate graphs for average path length



# Takeaways

- Difficult to test theories due to sheer size of applications needed and unavailability of testing frameworks / networks
- Systems programming is hard

**Thank you!**

